The frontier today: two labs define the edge

Both shipped new flagships in summer 2026. Capabilities and pricing are public — parameter counts are not.

OPENAI

GPT-5.6

Released July 9, 2026 · three variants: Luna (fast) · Terra (balanced) · Sol (most capable)

- Aimed at enterprise work, coding, scientific research, and cyberdefense — billed by OpenAI as its "strongest cybersecurity model yet"
- Sol: 54% more token-efficient on coding tasks than prior models (OpenAI claim)
- Emphasis on design judgment and safeguards that hold under "real-world adversarial pressure"

PARAMETERS: NOT DISCLOSED

ANTHROPIC

Claude Fable 5

Released June 9, 2026 · first "Mythos-class" model · same underlying model as restricted-access Claude Mythos 5

- Fable 5 = broad availability with new high-risk-area safeguards; Mythos 5 = trusted cyberdefense/infrastructure partners, fewer safeguards
- Anthropic: >10% above Claude Opus 4.8 on some benchmarks; strong software-engineering and knowledge-work results
- Pricing: $10 / $50 per million tokens (input / output)

PARAMETERS: NOT DISCLOSED

Name check: the voice transcript's "Fable" is not a recognition error — Claude Fable 5 is the confirmed Anthropic model name. Google (Gemini), xAI and others also ship frontier models; per scope, this draft compares the two labs above.

Sources: OpenAI (Jul 9, 2026); TechCrunch; CNBC (Jun 9 & Jul 8, 2026); Anthropic announcements. Details and links in speaker notes.

The scale trajectory: ≈30× since ChatGPT (est.)

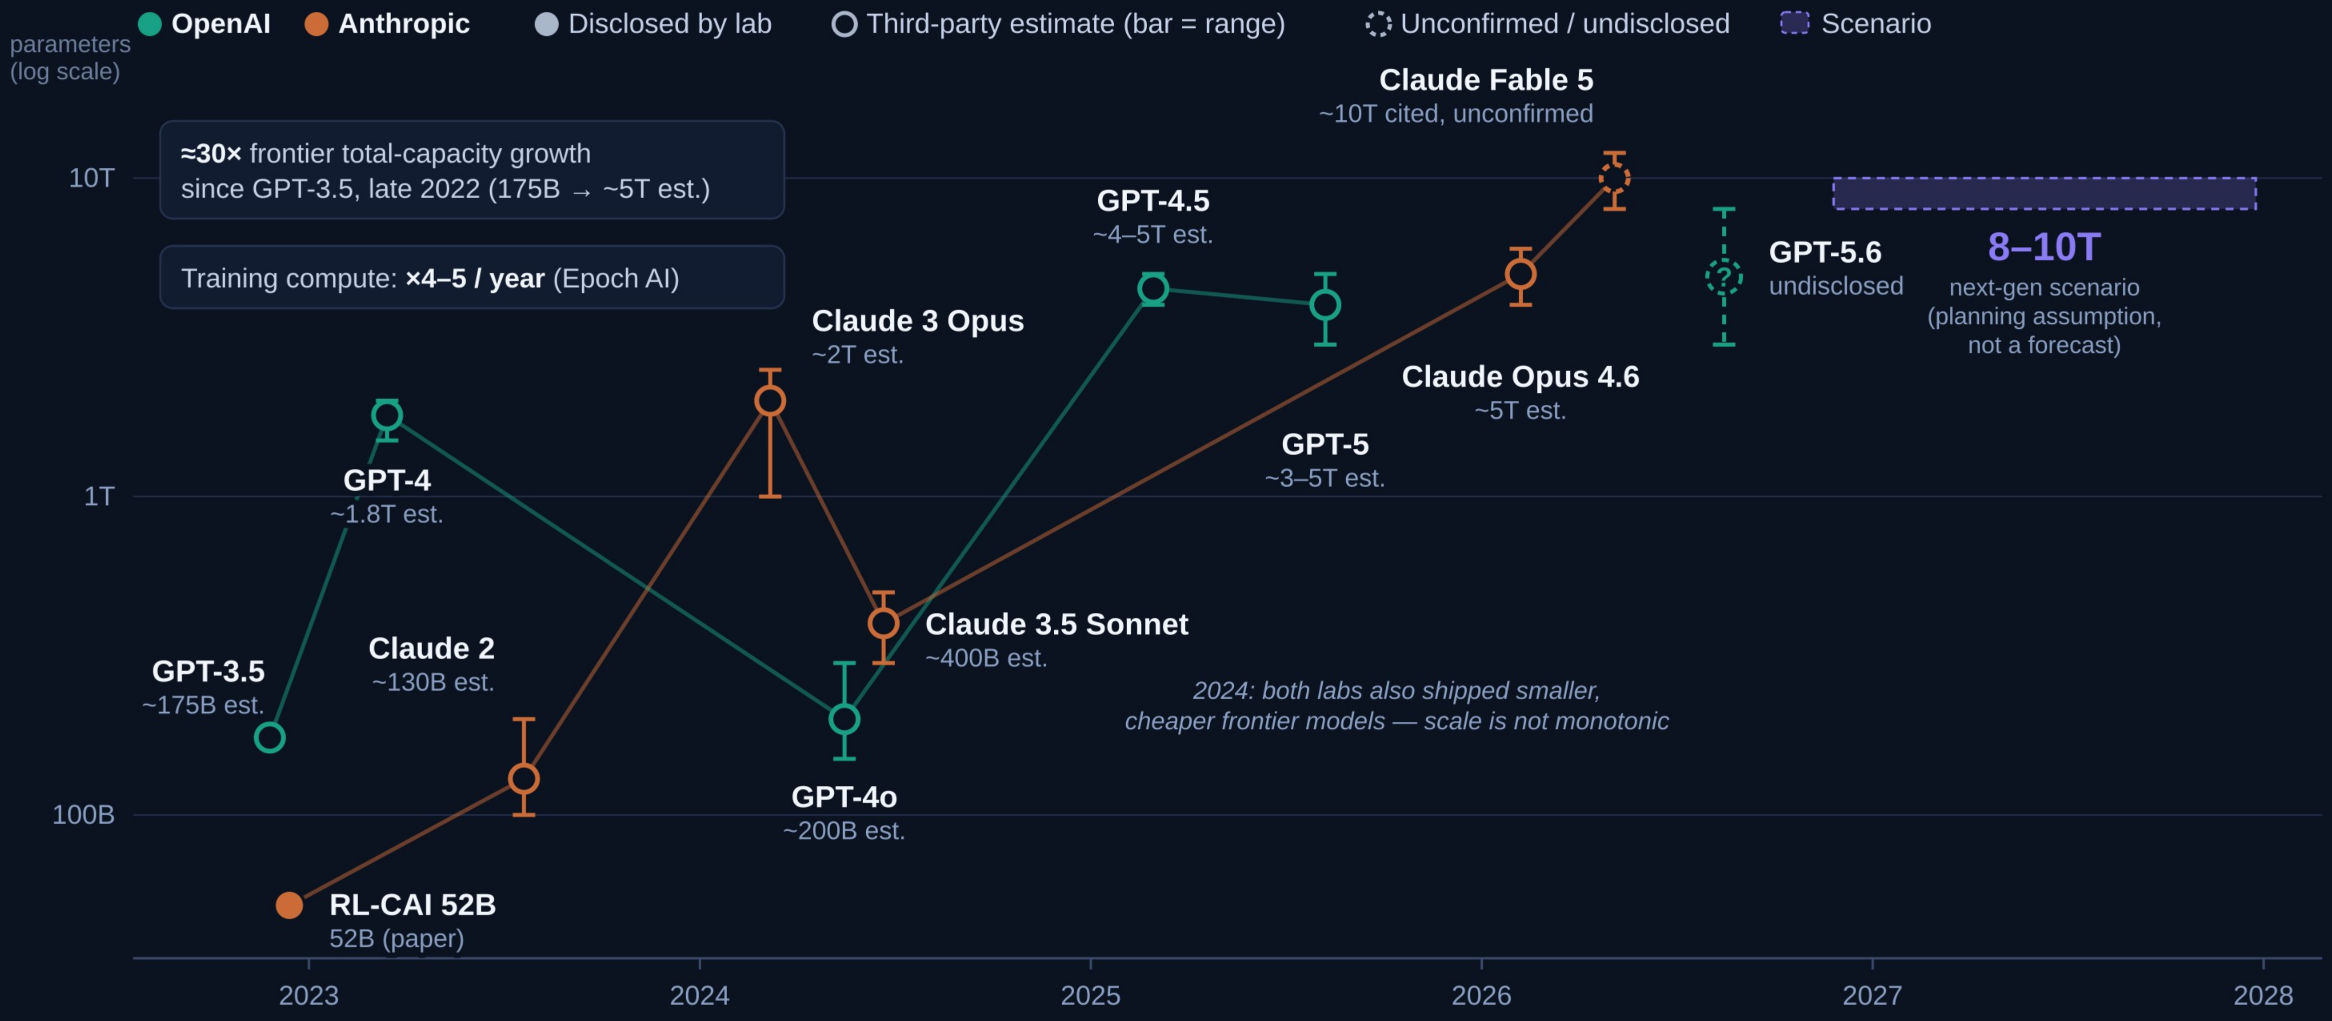

Reading this honestly: OpenAI last disclosed a flagship count with GPT-3 (2020); Anthropic never has for Claude flagships. Filled dots = lab-disclosed; hollow dots = third-party estimates (bars = published ranges). MoE splits “total” vs “active”: GPT-4 est. ~1.8T total, ~280B active per token. Log scale: each gridline is 10×. Sources: lab papers (GPT-3 lineage for GPT-3.5; RL-CAI 52B); SemiAnalysis (GPT-4); Epoch AI (GPT-4o, 3.5 Sonnet, compute trend); LifeArchitect (Claude 3 Opus); press/analyst reports (later flagships).

Next generation: the 8–10T scenario

SCENARIO — NOT A FORECAST

Working assumption for planning: next-generation frontier models reach roughly 8–10 trillion total parameters and deliver another major capability step.

8–10T

total parameters, next-gen frontier (planning assumption)

≈ 2× today's estimated ~5T frontier — one more historical scaling step, not a discontinuity

If the unconfirmed ~10T Mythos-class figures are right, this scale may already exist — making the scenario conservative on size

What supports it — verified trends

- Frontier training compute has grown ×4–5 per year for a decade (Epoch AI)
- Epoch judges ×4/yr feasible through 2030 — needing >5 GW training power
- 2026 frontier runs: 1e26–1e27 FLOPs, $200–500M per run — and still scaling
- Est. frontier size roughly doubled per generation since 2023 (~1.8T → ~5T), putting 8–10T on-trend for the next step

What stays uncertain

- No lab discloses counts — 8–10T cannot be verified against ground truth
- Scale is not monotonic: 2024's frontier shrank 5–10× before re-scaling
- Capability now leans on post-training RL and test-time compute — a 10T model is not automatically a leap
- MoE blurs "size": total vs active parameters differ ~5–10×

Recommended framing: plan capacity against the 8–10T scenario; treat "another major capability leap" as plausible but unproven — and track capability benchmarks, not parameter rumors, as the leading indicator.

Sources: Epoch AI (compute-trend and power analyses); scenario sizing: this team's working assumption, benchmarked against third-party estimates shown on slide 2.